

AICUMEN

The Acharya Method

Understanding AI

Class 6 — Part 2: Artificial Intelligence Curriculum

Companion to the AICUMEN Class 6 Computational Thinking Bank (Part 1)

Built for KRUU mentors

Mentor-led · Screen-light · Teaching AI, not just using it

PropelUp AI × KRUU

About This Book

This is the AI half of the AICUMEN Class 6 curriculum. Where Part 1 (Computational Thinking) sharpens the reasoning that sits underneath AI, Part 2 teaches AI itself: what it is, how it learns, how to read data, and how to use it responsibly. Nothing here asks a Class 6 student to write code or train a model. Every idea is taught through discussion, story, and unplugged reasoning first, matching how CBSE has designed AI literacy for this age group and how KRUU already teaches, through a mentor, not a machine.

The goal is a student who can explain what AI is doing and why, not just operate it.

Structure & Scope

Twenty hours across four modules, matching the CBSE Class 6 AI syllabus allocation exactly. Each module opens with a short concept explainer for the mentor, followed by four to five classroom activities in the same PROBLEM / RUN IT / EXTEND format used in the Part 1 problem bank, so the two halves of the curriculum feel like one book.

Every activity also carries its own SOCRATIC LADDER, a numbered sequence of clues built specifically for that question, each one a nudge rather than an answer. This is the piece meant for the app: a student attempts the problem, and only if they're stuck does Clue 1 appear, then Clue 2, then Clue 3, with the Reveal held back until the very end. The mentor (or the app, standing in for one) never states the answer first, the student always says it.

Module	AICUMEN Focus	CBSE Class 6 AI Learning Outcome (mapped)
Module 1	What Is AI?	Summarise AI concepts; describe AI vs. automation; differentiate supervised, unsupervised, reinforcement learning
Module 2	Basic Data Concepts	Organise and represent data across text, numbers, images, sound
Module 3	Pattern Recognition & Decision-Making	Recognise patterns in data; make decisions from observation
Module 4	Ethics & Digital Responsibility	Digital footprints, privacy, password safety, responsible use

A Note for Mentors

CBSE's own model anchors AI content with a dedicated Computer teacher. KRUU classrooms run differently, one mentor, one group, no separate Computer-teacher role. This book is written for that reality: every activity can be run by a single generalist mentor with paper, a whiteboard, and conversation. Where a digital tool would genuinely help (Module 2's data tables, for instance), it is offered as an optional RUN IT variation, never a requirement.

All scores, prices, names, and data used in examples are simulated for teaching purposes and are flagged as such where it matters.

MODULE 1

What Is AI?

5 hours

Machine intelligence vs. human intelligence, AI vs. automation, and the three ways machines learn.

Start here before any of the activities: a calculator follows fixed rules someone wrote in advance, that is automation. An AI system looks at examples and works out its own rules, that is learning. Once students can tell those two apart, the three learning types (supervised, unsupervised, reinforcement) are just three different ways of saying "how did it learn its rules."

1. AI or Just a Machine?

Everyday Tech

★☆☆☆☆

PROBLEM

Four gadgets are named: a kitchen weighing scale, a UPI app that flags a suspicious payment, a ceiling fan regulator, and a phone camera that recognises faces to unlock. Which of these are AI, and which are simple automation?

RUN IT

Hold up (or draw) four cards, one per gadget. Students physically move to an "AI" corner or an "Automation" corner of the room for each one, then defend their choice.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → For each gadget, ask: does it follow one instruction someone wrote in advance, or does it seem to have looked at many examples first?

Clue 2 → Look at the kitchen scale and the fan regulator together. What do the two of them have in common in how they work?

Clue 3 → Now look at the UPI fraud flag and the face-unlock camera. What would each of them have needed to see many times before, in order to work at all?

Reveal, if still stuck after Clue 3: Automation: the kitchen scale and fan regulator, both follow one fixed rule with no learning involved. AI: the UPI fraud flag and face-unlock camera, both learned a pattern from many past examples.

EXTEND

Ask: could the fan regulator ever become AI? What would have to change? (It would need to learn a pattern, for instance adjusting itself based on how hot the room usually gets at that time of day.)

2. Sorting Mangoes with Meera Aunty

Kirana / Market

★★☆☆☆

PROBLEM

Meera Aunty has sorted 40 mangoes into "ripe" and "raw" piles, and every mango already has a

little sticker saying which pile it belongs to. Her nephew watches all 40, then picks up mango number 41, which has no sticker, and predicts which pile it belongs to. What kind of learning is this?

RUN IT

Bring (or draw) 10 mango cards, 5 labelled "ripe" and 5 labelled "raw," with colour and firmness noted on each. Add one unlabelled card and have students predict its label using the pattern from the first 10.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → *What's the difference between having a labelled example and not having one?*

Clue 2 → *Did the first 40 mangoes already have their correct answer, ripe or raw, written on them?*

Clue 3 → *So is the nephew guessing mango 41's label from nothing at all, or from the pattern in 40 labelled examples?*

Reveal, if still stuck after Clue 3: This is supervised learning: the nephew learns from mangoes that already carry the correct label, then applies that pattern to a new, unlabelled case.

EXTEND

What would happen if half the stickers were wrong? Would the nephew still learn the right pattern? (This is a gentle first brush with the idea of noisy or incorrect training data.)

3. The Unlabelled Spice Box

Kirana / Market

★★★★☆

PROBLEM

A new shop assistant is handed a box with 30 loose spice packets, no labels, all mixed together. He is told to "group the similar ones together" without being told what the groups should be called. He ends up with piles by smell and colour: reddish-orange powders in one pile, dark whole pods in another, pale seeds in a third. What kind of learning does this match?

RUN IT

Give students a mixed set of 12–15 small objects (buttons, beads, pulses) with no labels and ask them to form their own groups, then explain the rule they used.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → *Were the spice packets given any names or labels to begin with?*

Clue 2 → *So how did the assistant decide which pile a packet belonged to, if nobody told him?*

Clue 3 → *If there's no correct answer given in advance, and the groups are discovered rather than told, what kind of learning does that sound like?*

Reveal, if still stuck after Clue 3: This is unsupervised learning: there are no labels at all, the assistant (or system) finds its own structure by comparing similarities.

EXTEND

Two students given the identical unlabelled box may form different groupings. Discuss: is one of them wrong, or can both groupings be valid?

4. Learning to Fly a Kite

Everyday Life

★★★★☆☆

PROBLEM

Aryan has never flown a kite before. Nobody gives him a rulebook. He tries pulling the string a little, the kite dips, he tries pulling harder, it climbs, he keeps adjusting based on what happens each time, and after twenty tries, he can keep the kite steady. Which of the three learning types does this match, and why?

RUN IT

Have a small group play a simple hand-clap rhythm game where a "reward" (a point, a clap-back) is given only for a correct pattern, so the group has to adjust their rhythm attempt by attempt.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → Did anyone hand Aryan a rulebook or a set of correct examples before he started flying the kite?

Clue 2 → What happens right after each pull of the string, does something tell him whether that pull was right or wrong?

Clue 3 → Is that "something" a labelled example, or more like a consequence, the kite rising or dipping, that he adjusts to?

Reveal, if still stuck after Clue 3: This is reinforcement learning: there's no example data at all, only trial, a consequence (the kite rising or falling), and adjustment toward a goal.

EXTEND

What is playing the role of the "reward" in Aryan's kite example? What would happen if the reward (the kite staying up) never came, no matter what he tried?

5. Name That Learning Type

Mixed Review

★★★★☆☆

PROBLEM

Four short scenarios are read aloud: (a) a music app groups songs into "playlists you might like" with no one telling it what the playlists should be, (b) a doctor's assistant predicts a disease from thousands of past patient records that were each marked with the correct diagnosis, (c) a robot vacuum learns the fastest route around a room only after bumping into furniture many times, (d) an app sorts old family photos into folders by comparing them to a set of pictures you already tagged with names. For each, name the learning type.

RUN IT

Read one scenario at a time. Students hold up one of three cards (Supervised / Unsupervised / Reinforcement) simultaneously, so answers cannot be copied from a neighbour.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → For each scenario, ask first: is there labelled example data given upfront, no labels at all, or just trial-and-error feedback?

Clue 2 → If there are no labels and no reward, but the system is grouping things by similarity, what does that sound like?

Clue 3 → If there IS a reward or consequence driving the behaviour, but no upfront correct-answer examples, what does that sound like instead?

Reveal, once all four are attempted: (a) Unsupervised, groups songs with no given labels. (b) Supervised, learns from records already labelled with the correct diagnosis. (c) Reinforcement, learns by trial, error, and consequence. (d) Supervised, learns from photos you already labelled with names.

EXTEND

Ask students to invent one new scenario of their own, from home or their school, for each of the three learning types.

MODULE 2

Basic Data Concepts

5 hours

Recognising data types, and organising raw information into something a machine, or a mentor, can actually read.

AI cannot learn from a mess. Before any pattern-spotting can happen, data has to be noticed, sorted into a type, and arranged. This module is entirely about that groundwork, using things students already handle every day: registers, receipts, recordings, and photographs.

1. What Kind of Data Is This?

School Life ★☆☆☆☆

PROBLEM

Four items are shown: the class attendance register, a voice note from the school announcement speaker, a photo of the school Republic Day parade, and a wedding invitation card. Sort each into one of four data types: text, number, image, or sound.

RUN IT

Bring four real or printed examples into class and have students physically sort them under four labelled headings on a table.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → Does this item contain written words, numbers, a picture you can see, or something you can only hear?

Clue 2 → Look closely at the attendance register, could it actually hold more than one data type at once?

Clue 3 → Sound is different from image, and text is different from number, which of the four fits each remaining item best?

Reveal, if still stuck after Clue 3: Attendance register: mainly numbers (with some text). Announcement voice note: sound. Republic Day photo: image. Wedding card: text.

EXTEND

Can one item hold more than one data type at once? (The attendance register has both text, names, and numbers, roll numbers and present/absent counts.)

2. Organise the Kirana Shelf

Kirana / Market ★★☆☆☆

PROBLEM

A shopkeeper has a messy list of 20 items scribbled on paper: rice, soap, toothpaste, dal, shampoo, wheat flour, and so on, with prices written next to each in no particular order. Organise this into a clean table with clear categories (for example, Groceries, Personal Care) and columns for item name and price.

RUN IT

Give groups the scrambled list on paper strips. They physically arrange the strips into a table on the floor or a large sheet before copying it into a notebook.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → Before organising anything, can you sort the 20 items into just two or three big groups?

Clue 2 → Within a group like "Personal Care," is there an order that would make the list easier to scan, like alphabetical?

Clue 3 → What information does each item need next to it to actually be useful, just a name, or a name and a price?

Reveal, if still stuck after Clue 3: A worked example table is provided in the Facilitator Guide, two category columns (Groceries, Personal Care) and a price column, alphabetised within each category.

EXTEND

Once organised, ask: which category has the most items? Which is more expensive on average? Notice how organising data makes new questions answerable that were invisible in the messy list.

3. Reading the Monsoon Chart

Weather / Geography ★★☆☆☆

PROBLEM

A simple bar chart shows rainfall (in mm, simulated figures) for June through September in a small town: June 180, July 310, August 260, September 140. Which month had the heaviest rainfall? How much more rain fell in July than September?

RUN IT

Draw the four bars on the board without numbers first. Have students estimate values purely by comparing bar heights before revealing the numbers.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → Which bar on the chart reaches the highest point?

Clue 2 → Now find July's bar and September's bar specifically, which one is taller, and by roughly how much?

Clue 3 → To find "how much more" rain fell in one month than another, what do you do with the two numbers, add them or subtract them?

Reveal, if still stuck after Clue 3: Heaviest rainfall: July (310 mm). Difference between July and September: $310 - 140 = 170$ mm. (Rainfall figures here are simulated for the exercise, not real recorded data for any specific town.)

EXTEND

If a farmer had to pick the single best month to plant a crop that dislikes too much water but needs some rain, which month from this chart looks safest, and why?

4. Build Your Own Data Table

Classroom Activity



PROBLEM

Each student notes down three things about five classmates: favourite subject, number of siblings, and preferred sport. This is scattered, spoken information. Turn it into a single organised table.

RUN IT

Physical interview activity: students walk around, ask five classmates the three questions, and fill in a blank table template as they go.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → *What three pieces of information are you collecting from each classmate?*

Clue 2 → *If those three pieces become your column headers, what should go in the very first column of all?*

Clue 3 → *Once every row is filled in, how could you find the class's most common favourite subject just by looking at the finished table?*

Guidance, not a locked answer: No single fixed answer, this activity produces real classroom data. Model the header row (Name / Subject / Siblings / Sport) before students begin, then count occurrences together as a class.

EXTEND

Once every student has a table, combine everyone's tables on the board into one class-wide table. What is the most common favourite subject in the whole class?

MODULE 3

Pattern Recognition & Decision-Making

5 hours

Spotting what repeats, then acting on it, the two-step move every AI system makes, done here entirely by hand.

This module bridges Module 1 (how machines learn) and Module 4 (what can go wrong). Recognising a pattern in data and then making a decision based on it is exactly what a supervised or reinforcement system does internally. Doing it by hand first makes the "black box" feel a lot less mysterious later.

1. The Dabbawala's Rhythm

Mumbai / Logistics

★★☆☆☆

PROBLEM

A dabbawala collects lunchboxes from five buildings in this order every single day: A, C, B, A, C, B, A, C, ? What comes next, and how did you know?

RUN IT

Use physical building cards (or students standing in a line) and have the class physically continue the sequence by stepping into the next position.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → Look at the sequence so far, does anything repeat after a fixed number of steps?

Clue 2 → Count how many buildings appear before the pattern starts over. What is that number?

Clue 3 → If the cycle is A, C, B and it has just repeated up to ...A, C, what comes right after C in that cycle?

Reveal, if still stuck after Clue 3: The sequence repeats in a fixed 3-step cycle: A, C, B. After ...A, C, B, A, C, the next is B.

EXTEND

If one day the dabbawala had to skip building C because it was closed, what pattern would the delivery sequence show that day?

2. Should the Library Say Yes or No?

School Life

★★★☆☆

PROBLEM

A school library has a simple rule for its digital system: if a book is more than 14 days overdue AND the student has more than 2 books already borrowed, flag the account. Given three students' records (Student 1: 10 days overdue, 1 book; Student 2: 20 days overdue, 3 books; Student 3: 16 days overdue, 1 book), which accounts get flagged?

RUN IT

Give each group a stack of "student record" cards with the two numbers on each, and have them

physically sort cards into "Flag" and "No Flag" piles by applying the rule.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → What are the two conditions that must BOTH be true for an account to be flagged?

Clue 2 → Check Student 1 against both conditions, does either one actually hold true?

Clue 3 → Now check Student 3, they're well past 14 days overdue, but what about the second condition, the number of books?

Reveal, if still stuck after Clue 3: Only Student 2 is flagged (20 days overdue AND 3 books, both conditions true). Students 1 and 3 each fail at least one condition. This is a rule-based decision, not a learned one, useful for contrasting with AI from Module 1.

EXTEND

This rule was written by a person, not learned from data. How is this different from an AI system deciding the same thing after studying thousands of past overdue records?

3. Predicting the Next Over

Cricket ★★★★★☆

PROBLEM

In a simulated T20 match, a batter has scored 4, 6, 4, 6, 4 runs off her last five balls faced from one bowler. Based only on this pattern, predict roughly what she might score off the next ball, and explain why a real prediction would need far more data than this.

RUN IT

Chart the five scores as dots on a simple number line or bar drawing, then have students mark where they think the sixth dot should sit and defend the choice.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → Look at the five scores, do they seem to alternate between two particular numbers?

Clue 2 → If that alternating pattern truly holds, what would come right after ...4, 6, 4?

Clue 3 → Now ask: is five balls really enough data to trust this prediction in a real match? What real information is missing?

Reveal, if still stuck after Clue 3: The visible pattern alternates high-low (4, 6, 4, 6, 4), so 6 is a reasonable pattern-based guess, but the real teaching point is that five balls is far too little data, real systems would want the bowler's full over-by-over history, delivery type, and field placement. (All figures simulated for the exercise.)

EXTEND

What other information (beyond these five balls) would a real AI system want before making this prediction? List at least three.

4. Spot the Odd One in the Data

Data Detective ★★★★★☆

PROBLEM

A shop's daily sales for one week (in rupees, simulated): 2,000 / 2,200 / 1,950 / 2,100 / 15,000 / 2,050 / 2,150. One number does not fit the pattern of the rest. Which one, and what might

explain it in real life?

RUN IT

Plot the seven values as bars on the board (heights roughly to scale), the mismatch becomes visually obvious before any calculation.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → *Look at six of the seven numbers, ignoring the one that looks different, are they roughly close to each other?*

Clue 2 → *Which single number breaks that closeness the most?*

Clue 3 → *In real life, what kind of day could cause one day's sales to jump far above all the rest?*

Reveal, if still stuck after Clue 3: Rs. 15,000 is the anomaly, roughly seven times the typical daily figure. Possible explanations: a festival rush day, a bulk order, or a data-entry error, the pattern alone cannot tell us which, only that it needs a closer look. (Figures simulated for the exercise.)

EXTEND

An AI fraud-detection system uses exactly this idea, spotting a transaction that does not match a person's usual pattern. What is one everyday example where you'd want a system to flag something unusual like this?

MODULE 4

Ethics & Digital Responsibility

5 hours

Privacy, fairness, and safe habits, so a student can use AI and the internet with judgement, not just enthusiasm.

This module has no right-or-wrong worked "solutions" in the usual sense, the aim is reasoned discussion, not a single correct number. Each activity closes with guiding questions for the mentor rather than a locked answer, and ties directly to India's Digital Personal Data Protection (DPDP) Act 2023, the governing privacy framework for this age group.

1. Whose Photo Is It, Really?

Digital Footprint

★ ★ ☆ ☆ ☆

PROBLEM

A student posts a photo from a school annual day function online, tagging four classmates by name, without asking them first. Is this a problem? If so, for whom?

RUN IT

Role-play in pairs: one student plays the poster, one plays a tagged classmate who was not asked. Act out the conversation that should have happened before posting.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → Was anyone in the photo actually asked before it was posted and tagged with their name?

Clue 2 → Does tagging someone add information about them to the internet, even though they didn't post it themselves?

Clue 3 → If a tagged classmate later wanted that photo taken down, could they always make that happen easily and quickly?

Discussion points to arrive at together: Consent matters even for photos taken in a shared space; tagging adds to someone else's digital footprint, not just the poster's; asking first is a small habit with a real, lasting difference.

EXTEND

What is a "digital footprint," and how is tagging someone else's name and face part of building their digital footprint, not just your own?

2. Fair or Unfair?

Bias in AI

★ ★ ★ ☆ ☆

PROBLEM

A school's new sports-recommendation app only ever suggests cricket to boys and dance to girls, no matter what a student actually enters as their interest. What has gone wrong here, and where would this problem most likely have come from?

RUN IT

Small groups design (on paper) what training examples they would feed the app instead, to avoid this exact problem, and present their fix to the class.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → Does the app's suggestion actually change based on what a student types, or does it seem to depend on something else entirely?

Clue 2 → If this app learned from past examples, like Module 1's supervised learning, where would those past examples have originally come from?

Clue 3 → If those past examples already reflected old assumptions about who plays what sport, what would the app end up learning to repeat?

Discussion points to arrive at together: This is a bias problem, most likely rooted in training data that reflected old assumptions (past enrolment records skewed by gender). The fix starts with better, more varied examples, not just a smarter algorithm.

EXTEND

Connect back to Module 1: if this app learned from supervised examples, what does that suggest about the examples it was trained on in the first place?

3. Password Power

Online Safety ★☆☆☆☆

PROBLEM

Compare four passwords: "123456," "ravi2015" (a birth year), "R@vi!2015#Cricket," and "correcthorsebatterystaple." Rank them from weakest to strongest and explain your ranking.

RUN IT

Write all four on cards and have students physically arrange them in a line from weakest to strongest, then debate any disagreements.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → Which of the four passwords could someone guess just by knowing your name and birth year?

Clue 2 → Does adding symbols to a personal detail, like a birth year, make it truly unpredictable, or just longer and fiddlier to type?

Clue 3 → What makes "correcthorsebatterystaple" hard to guess, even though it has no symbols or numbers at all?

Reveal, if still stuck after Clue 3: Weakest to strongest, roughly: "123456" (an extremely common, guessable sequence), "ravi2015" (personal and predictable), "R@vi!2015#Cricket" (longer and mixed, but still built from guessable personal details), "correcthorsebatterystaple" (long and unpredictable, hard to guess by a pattern-matching attacker).

EXTEND

Why might a long, unusual phrase sometimes be just as strong as, or stronger than, a shorter password stuffed with symbols?

4. The Right to Say No

Privacy / DPDP



PROBLEM

A free learning app asks for permission to access a student's photo gallery and contact list before letting them start a simple maths quiz. Does the app actually need this information to do its job? What should a student, or a parent, do here?

RUN IT

Have students list, for three apps they actually use, what permission each one asks for versus what it actually needs to do its stated job.

SOCRATIC LADDER • reveal one clue per student attempt

Clue 1 → *What does a simple maths quiz actually need to know about you in order to function?*

Clue 2 → *Does "photo gallery" or "contact list" appear anywhere in that list of real needs?*

Clue 3 → *India's DPDP Act says consent should be specific and informed, not a single blanket click. What does an unrelated permission request suggest about that standard?*

Discussion points to arrive at together: Consent should be tied to what a service actually needs; unrelated permission requests are a signal to pause and ask why; it is reasonable, and encouraged, to say no or check with an adult first.

EXTEND

India's DPDP Act 2023 gives individuals a right to meaningful consent, permission that is specific and informed, not just a single blanket "Accept All" click. Why does a maths-quiz app asking for contacts fail that standard?

Teaching Through Questions: The Acharya Script

AICUMEN is built on the Acharya Method: a student reasons their way to an answer, the mentor guides with questions, in the spirit of the guru-shishya tradition. This section gives mentors a repeatable script for the moment that comes up most often in Module 1: a student asking outright, "Is this supervised or unsupervised?" The same ladder works for almost any concept question in this book.

The Four-Rung Ladder

Climb one rung at a time. Move to the next rung only after a genuine pause, at least 10 to 15 seconds of real think time, not an immediate jump to giving the answer.

Rung 1 — Reflect: turn the question back. "What do you think, and why?"

Rung 2 — Anchor: point back to something the student already knows or already sorted earlier in the activity.

Rung 3 — Narrow: ask one comparison question that isolates the one detail that decides the answer.

Rung 4 — Frame: offer a half-built sentence for the student to complete. Never hand over the finished answer yourself, let the student say the word first.

Worked Script: "Is this supervised or unsupervised?"

Setting: a student is stuck partway through Module 1, Activity 2 ("Sorting Mangoes with Meera Aunty").

Who	What They Say
Student	<i>Ma'am, is this supervised or unsupervised learning?</i>
Mentor (Rung 1 — Reflect)	Good question. Before I answer, tell me, what's different between the two, in your own words?
Student	<i>Um... one has labels and one doesn't?</i>
Mentor (Rung 2 — Anchor)	Exactly right so far. Now look back at the mango problem. Did the first 40 mangoes have anything written on them?
Student	<i>Yes, stickers saying ripe or raw.</i>
Mentor (Rung 3 — Narrow)	So those 40 mangoes already came with the correct answer attached. What does that tell you about mango number 41, the one with no sticker?
Student	<i>He's trying to guess its label using the pattern from the other ones.</i>
Mentor (Rung 4 — Frame (only if still unsure))	So is this a case where the learner sees examples that already have the correct answer marked, or a case where it has to find its own groups with no answers given at all?
Student	<i>The first one. Supervised!</i>
Mentor	Exactly. Now tell me why, in one full sentence, so it sticks.

The one rule that matters most

The mentor never says the word "supervised" or "unsupervised" first. The mentor only confirms it once the student has said it. If a rung is skipped and the mentor names the answer before the student does, the exercise has quietly turned back into a lecture.

Quick Reference: Other Questions That Come Up Often

The same four-rung habit applies well beyond the supervised/unsupervised question. A few more starting points:

If a Student Asks...	Don't Just Say	Ask Instead
"Is this AI or automation?"	<i>"It's automation, because..."</i>	"Does it follow one fixed rule every time, or does it seem to change its behaviour after seeing examples?"
"Is this reinforcement learning?"	<i>"Yes, because there's a reward."</i>	"Is anyone giving correct-answer examples here, or is the learner just trying things and seeing what happens?"
"Is this fair or biased?"	<i>"Yes, it's biased because..."</i>	"Where do you think the app's assumptions came from in the first place?"
"What's the pattern here?"	<i>States the rule outright.</i>	"What changes as you move from one term to the next? Is that change the same every time?"

When to Stop Asking and Just Tell

The Acharya Method is a tool, not a rule to enforce at any cost. If a student has genuinely tried across two or three rungs and is getting frustrated or anxious rather than curious, or if time is short, it is completely fine to give the answer plainly and move on. The point of the ladder is better understanding and a confident student, not a struggle for its own sake. Watch the student, not just the ladder.

Facilitator Guide

Notes for the KRUU mentor delivering this book, since Part 2 asks for a different posture than Part 1.

Running AI Content Without a Computer Teacher

CBSE's own middle-stage model anchors AI Literacy content with a dedicated Computer teacher. KRUU's mentor-led structure does not have that role, and this book is built so it does not need one. Every activity above runs on paper, cards, or spoken scenario, with a digital tool offered only as an optional RUN IT variation. If a school does have computer access, Module 2's data-table activities are the natural place to bring in a spreadsheet, but it is never required to complete the module.

Sequencing Across the Term

Modules 1 and 2 build the vocabulary a student needs before Module 3 makes sense, teach them in order. Module 4 (Ethics) can be split and threaded across the other three rather than saved for the very end, a five-minute ethics discussion after a Module 1 or 3 activity often lands better than a standalone 5-hour block.

Assessment, In Keeping With CBSE Guidance

CBSE is explicit that Middle Stage AI assessment should move away from rote testing. This book supports three formats already built into the activities above:

- Recognition: the "Name That Learning Type" and "AI or Just a Machine?" style of quick-fire classification.
- Application: activities like "Sorting Mangoes" and "Should the Library Say Yes or No?" that ask students to apply a concept to a new case.
- Reflection: every Module 4 activity, and the "Extend" prompt in most others, which has no single correct answer and rewards reasoning.

No activity in this book requires "write a program" or memorised definitions, matching CBSE's own stated intent for this age group.

A Worked Table for "Organise the Kirana Shelf" (Module 2, Activity 2)

A sample completed table a mentor can use to check student work:

Category	Item	Price (Rs., simulated)
Groceries	Rice (1 kg)	65
Groceries	Wheat flour (1 kg)	48
Groceries	Dal (1 kg)	110
Personal Care	Soap	35
Personal Care	Toothpaste	55
Personal Care	Shampoo (small)	90

(Prices above are simulated for teaching purposes.)

What Comes Next

This book covers the Class 6 AI syllabus in full (20 hours, four modules). Classes 7 and 8 build on these same foundations with AI domains (Computer Vision, NLP, Data Science), predictive techniques (regression, classification, clustering), and the AI project lifecycle, and can follow this same house style when the team is ready to build them.